

# The Spark Foundation

Name : Priya Keshav Patil

## Data Science and Business Analytics

### Task 1: Prediction using Supervised ML

This Task is to predict the percentage of a student based on the no. of study hours. This is a simple linear regression task as it involves just two variables.

```
In [14]: #importing the libraries  
import pandas as pd  
import numpy as np  
import matplotlib.pyplot as plt  
import seaborn as sns
```

```
In [3]: #read the dataset  
df=pd.read_csv("dataset.csv")  
df
```

```
Out[3]:
```

	Hours	Scores
0	2.5	21
1	5.1	47
2	3.2	27
3	8.5	75
4	3.5	30
5	1.5	20
6	9.2	88
7	5.5	60
8	8.3	81
9	2.7	25
10	7.7	85
11	5.9	62
12	4.5	41
13	3.3	42
14	1.1	17
15	8.9	95
16	2.5	30
17	1.9	24
18	6.1	67

	Hours	Scores
19	7.4	69
20	2.7	30
21	4.8	54
22	3.8	35
23	6.9	76
24	7.8	86

In [4]: `df.shape`

Out[4]: (25, 2)

In [6]: `df.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 25 entries, 0 to 24
Data columns (total 2 columns):
 #   Column  Non-Null Count  Dtype  
---  -
 0   Hours   25 non-null     float64
 1   Scores  25 non-null     int64   
dtypes: float64(1), int64(1)
memory usage: 464.0 bytes
```

In [7]: `df.describe()`


Out[7]:

	Hours	Scores
count	25.000000	25.000000
mean	5.012000	51.480000
std	2.525094	25.286887
min	1.100000	17.000000
25%	2.700000	30.000000
50%	4.800000	47.000000
75%	7.400000	75.000000
max	9.200000	95.000000

## Exploring the dataset

In [12]:

```
plt.scatter(df['Hours'],df['Scores'],color="red",marker="o")
plt.title("Hours Vs Scores")
plt.xlabel("Study Hours")
plt.ylabel("percentage scored")
plt.show()
```




```
In [13]: df.corr()
```

```
Out[13]:
```

	Hours	Scores
Hours	1.000000	0.976191
Scores	0.976191	1.000000

```
In [15]: sns.lmplot(x="Hours",y="Scores",data=df)
plt.title("plotting the regression line")
```

```
Out[15]: Text(0.5, 1.0, 'plotting the regression line')
```



From the graph above, we can say that with the increase of Study Hours(x) there is an increase in the Scores obtained(y)

## Extracting the independent and dependent variable

```
In [18]: x=df.iloc[:, :-1].values  
         y=df.iloc[:, -1].values
```

```
Out[18]: (array([[2.5],  
                [5.1],  
                [3.2],  
                [8.5],  
                [3.5],  
                [1.5],  
                [9.2],  
                [5.5],  
                [8.3],  
                [2.7],  
                [7.7],  
                [5.9],  
                [4.5],  
                [3.3],  
                [1.1],  
                [8.9],  
                [2.5],  
                [1.9],  
                [6.1],  
                [7.4],  
                [2.7],  
                [4.8],  
                [3.8],  
                [6.9],  
                [7.8]]),  
         array([21, 47, 27, 75, 30, 20, 88, 60, 81, 25, 85, 62, 41, 42, 17, 95, 30,  
                24, 67, 69, 30, 54, 35, 76, 86], dtype=int64))
```

## Splitting dataset into training and testing set

```
In [19]: from sklearn.model_selection import train_test_split  
         x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=1/3,random_state=0)
```

## Training the Simple Linear Regression Model on the Training set

```
In [20]: from sklearn.linear_model import LinearRegression  
         regressor=LinearRegression()  
         regressor.fit(x_train,y_train)
```

```
Out[20]: LinearRegression()
```

## Predicting the Test set results

```
In [22]: y_pred=regressor.predict(x_test)  
         y_pred
```

```
Out[22]: array([17.04289179, 33.51695377, 74.21757747, 26.73351648, 59.68164043,  
                39.33132858, 20.91914167, 78.09382734, 69.37226512])
```

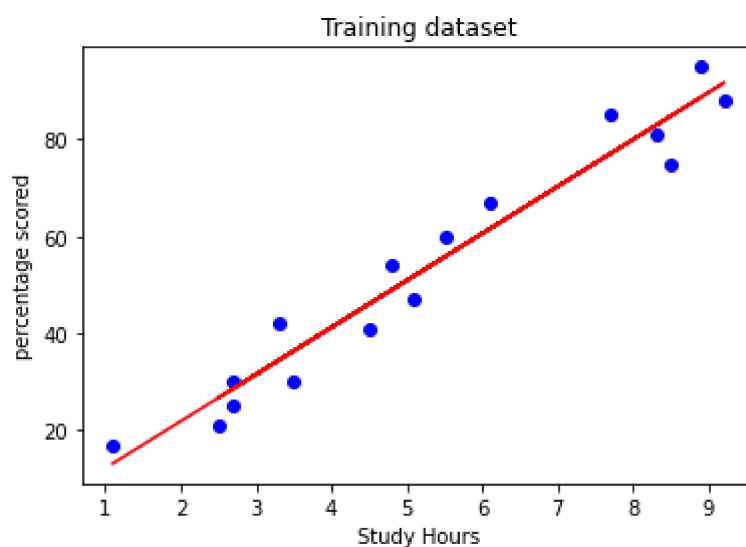
```
In [24]: # comparison between actual and predicted Score  
         df1=pd.DataFrame({'Actual': y_test, 'predicted':y_pred})  
         df1
```

Out[24]:

	Actual	predicted
0	20	17.042892
1	27	33.516954
2	69	74.217577
3	30	26.733516
4	62	59.681640
5	35	39.331329
6	24	20.919142
7	86	78.093827
8	76	69.372265


## Visualizing the Training set results

```
In [25]: plt.scatter(x_train,y_train,color='blue')
plt.plot(x_train,regressor.predict(x_train),color='red')
plt.title("Training dataset")
plt.xlabel("Study Hours")
plt.ylabel("percentage scored")
plt.show()
```




## Visualizing the Test set results

```
In [26]: plt.scatter(x_test,y_test,color='blue')
plt.plot(x_train,regressor.predict(x_train),color='red')
plt.title("Testing dataset")
plt.xlabel("Study Hours")
plt.ylabel("percentage scored")
plt.show()
```




```
In [27]: # checking for the correlations
plt.figure(figsize=(10,5))
sns.heatmap(df.corr(),annot=True,cmap="YlOrBr",annot_kws={'fontsize':12})
plt.xticks(fontsize=12)
plt.yticks(fontsize=12)
plt.show()
```



## Visualizing the Comparison between actual and predicted scores

```
In [32]: plt.scatter(y_test,y_pred,c='r')
plt.plot(y_test,y_pred,c='b')
plt.title("Score Vs Predicted Score")
plt.xlabel("Score")
plt.ylabel("Predicted Score")
```

```
Out[32]: Text(0, 0.5, 'Predicted Score')
```



What will be predicted score if a student studies for 9.25 hrs/day?

prediction through our model

```
In [33]: Hours=np.array([[9.25]])
predict=regressor.predict(Hours)
print("No of Hours:",format(Hours))
print("predicted Score:",format(predict[0]))
```

```
No of Hours: [[9.25]]
predicted Score: 92.14523314523314
```

Checking accuracy of our model

```
In [34]: print("train:",regressor.score(x_train,y_train)*100)
print("test:",regressor.score(x_test,y_test)*100)
```

```
train: 95.01107277744313
test: 95.5570080138813
```

Finding mean absolute error,  $r^2$  score error and mean squared error

```
In [36]: from sklearn import metrics
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error

print("Mean absolute error:",metrics.mean_absolute_error(y_test,regressor.predict(x_
print("r^2 score error:",r2_score(y_test,regressor.predict(x_test)))
print("Mean square error:",mean_squared_error(y_test,regressor.predict(x_test)))
```

```
Mean absolute error: 4.691397441397438
r^2 score error: 0.955570080138813
Mean square error: 25.463280738222547
```

Mean absolute error: 4.69137441397438 which is quite accurate model for predicting the result